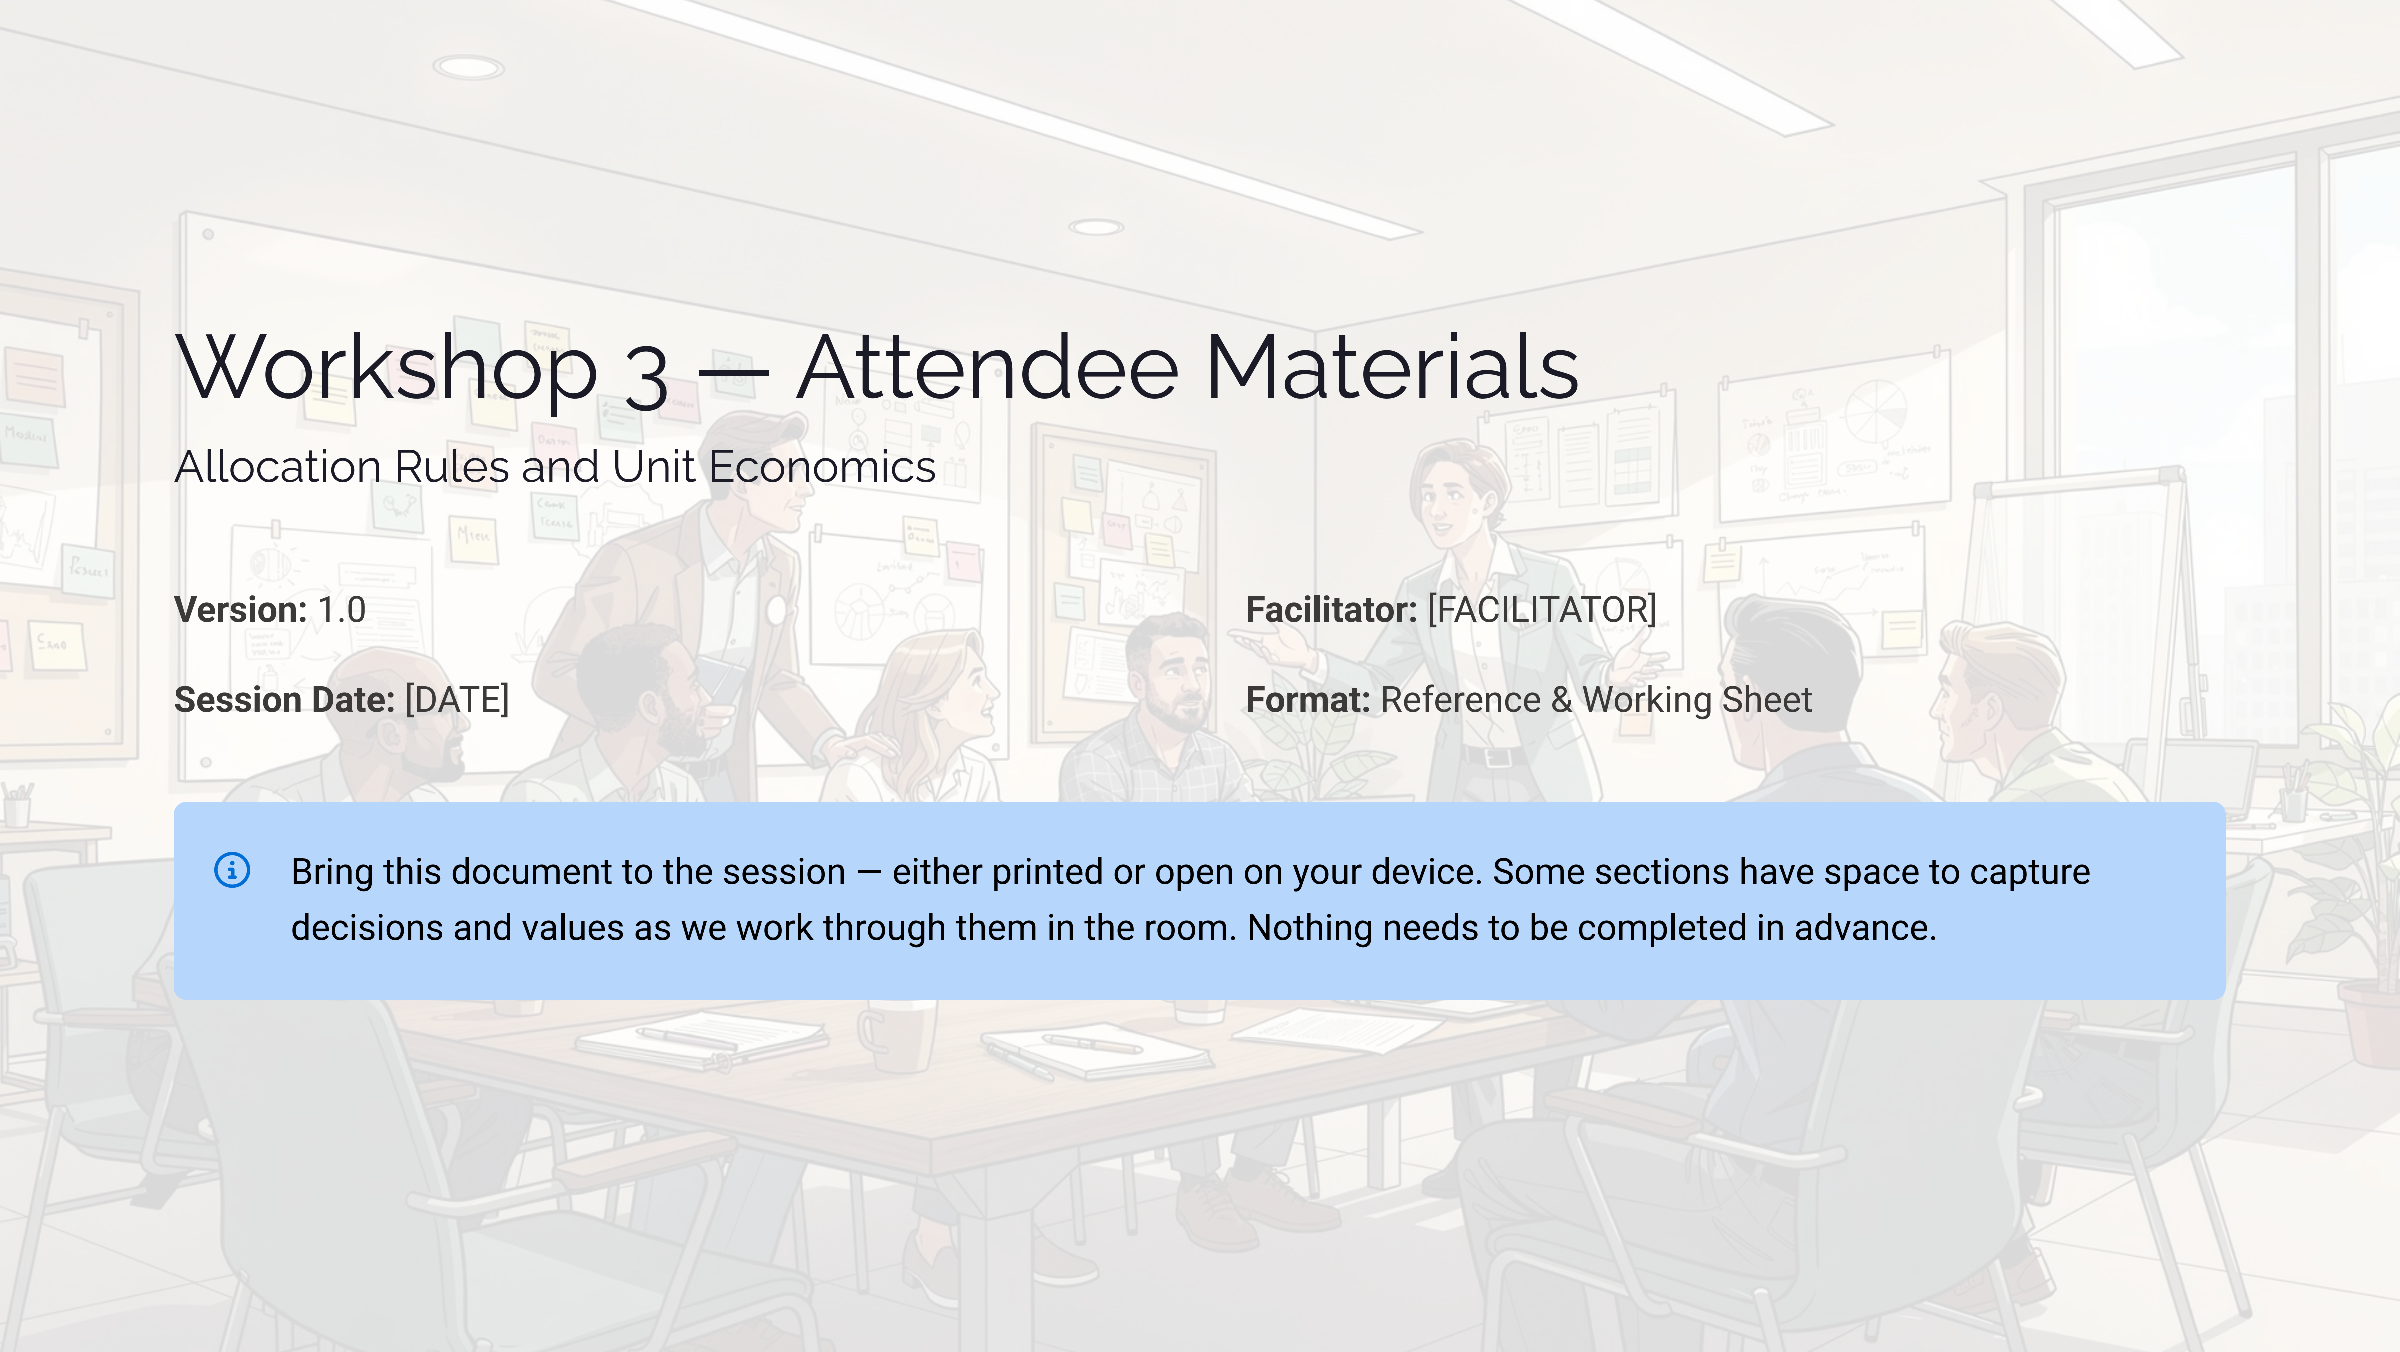

Workshop 3 — Attendee Materials

Allocation Rules and Unit Economics

Version: 1.0

Session Date: [DATE]

Facilitator: [FACILITATOR]

Format: Reference & Working Sheet

> Bring this document to the session — either printed or open on your device. Some sections have space to capture decisions and values as we work through them in the room. Nothing needs to be completed in advance.

What We Are Doing Today

Workshop 3 is where the cost model becomes useful to leadership. We are moving from **total cost** to **unit economics** — cost per useful outcome.

Allocation Rules

Agree allocation rules for all shared costs across the platform

Unit Metrics

Establish baseline values for our three to five unit metrics

Showback View

First-pass mapping of cost to teams or namespaces

Clear Narrative

Connect platform cost directly to business outcomes

Why This Matters

Total spend is a number. Unit economics is what that number *means*.

> Instead of "we spend X on Kubernetes", we want to be able to say something leadership and Finance can actually act on.

Production Application Cost

"A production application costs **Y per month**."

Tenant Cost

"Each tenant costs **Z per quarter**."

Utilisation Gap

"We pay for capacity but only use **A% of it** effectively."

These are the numbers leadership and Finance can act on — not raw spend totals.

Section 1 — Allocation Rules

We will agree on these together in the session. Record what is decided in each area below.

Shared Platform Costs

Option	Description	Selected?
A	Proportional by usage	
B	Equal split across platforms	
C	Shared unallocated bucket	

Decision

Reason

Platform Tooling Allocation

Option	Description	Selected?
A	Allocate by cluster or node count per platform	
B	Allocate equally across platforms	
C	Hold as shared — do not allocate	

Decision

Reason

Operating Effort Allocation

Option	Description	Selected?
A	Allocate by estimated time per platform	
B	Allocate equally across platforms	
C	Hold as platform-level cost — do not allocate to workloads	

Decision

Reason

Minimum Tagging Standard

Consistent tagging and labelling is the foundation of accurate cost allocation. Capture the current state and agreed minimum standard below.

Current Tags and Labels in Use

Tag / Label	What It Identifies	Coverage

Minimum Standard Agreed

Gaps to Close

⚠ Tagging gaps will directly limit the accuracy of your unit metrics. Prioritise closing these gaps before Workshop 4.

Section 2 — Unit Metrics

Our shortlist from Workshop 1. We will select three to five metrics and build baseline values for each during this session.

Cost per Production App / Month

Total platform cost divided by number of production applications running

Cost per Tenant or Namespace / Month

Total cost divided by number of active tenants or Kubernetes namespaces

Cost per Environment

Prod vs non-prod split — understanding the cost of lower environments

Effective Cost per vCPU-Hour Used

Compute cost adjusted for actual utilisation, not just provisioned capacity

Ops Cost per Cluster / Month

Operating effort cost allocated per cluster under management

Other

Additional metric identified by the group during the session:

____

Metric Worksheets — Metrics 1 & 2

Complete these worksheets together during the session. Record numerator, denominator, data sources, baseline value, and confidence rating for each metric.

Metric 1: [METRIC NAME]

Field	Value
Numerator (costs included)	
Denominator (what we divide by)	
Data sources	
Baseline value	
Confidence rating	H / M / L

What this number means:

What it enables:

Metric 2: [METRIC NAME]

Field	Value
Numerator (costs included)	
Denominator (what we divide by)	
Data sources	
Baseline value	
Confidence rating	H / M / L

What this number means:

What it enables:

Metric 3: [METRIC NAME]

Field	Value
Numerator (costs included)	
Denominator (what we divide by)	
Data sources	
Baseline value	
Confidence rating	H / M / L

What this number means:

What it enables:

Metric 4: [METRIC NAME]

Field	Value
Numerator (costs included)	
Denominator (what we divide by)	
Data sources	
Baseline value	
Confidence rating	H / M / L

What this number means:

What it enables:

Metric 5: [METRIC NAME]

Field	Value
Numerator (costs included)	
Denominator (what we divide by)	
Data sources	
Baseline value	
Confidence rating	H / M / L

What this number means:

What it enables:

Section 3 — Unit Economics Summary

Complete this together at the end of the metric-building exercise. This becomes the single reference table for the session output.

Metric	Baseline Value	Confidence	Key Finding
[METRIC 1]		H / M / L	
[METRIC 2]		H / M / L	
[METRIC 3]		H / M / L	
[METRIC 4]		H / M / L	
[METRIC 5]		H / M / L	

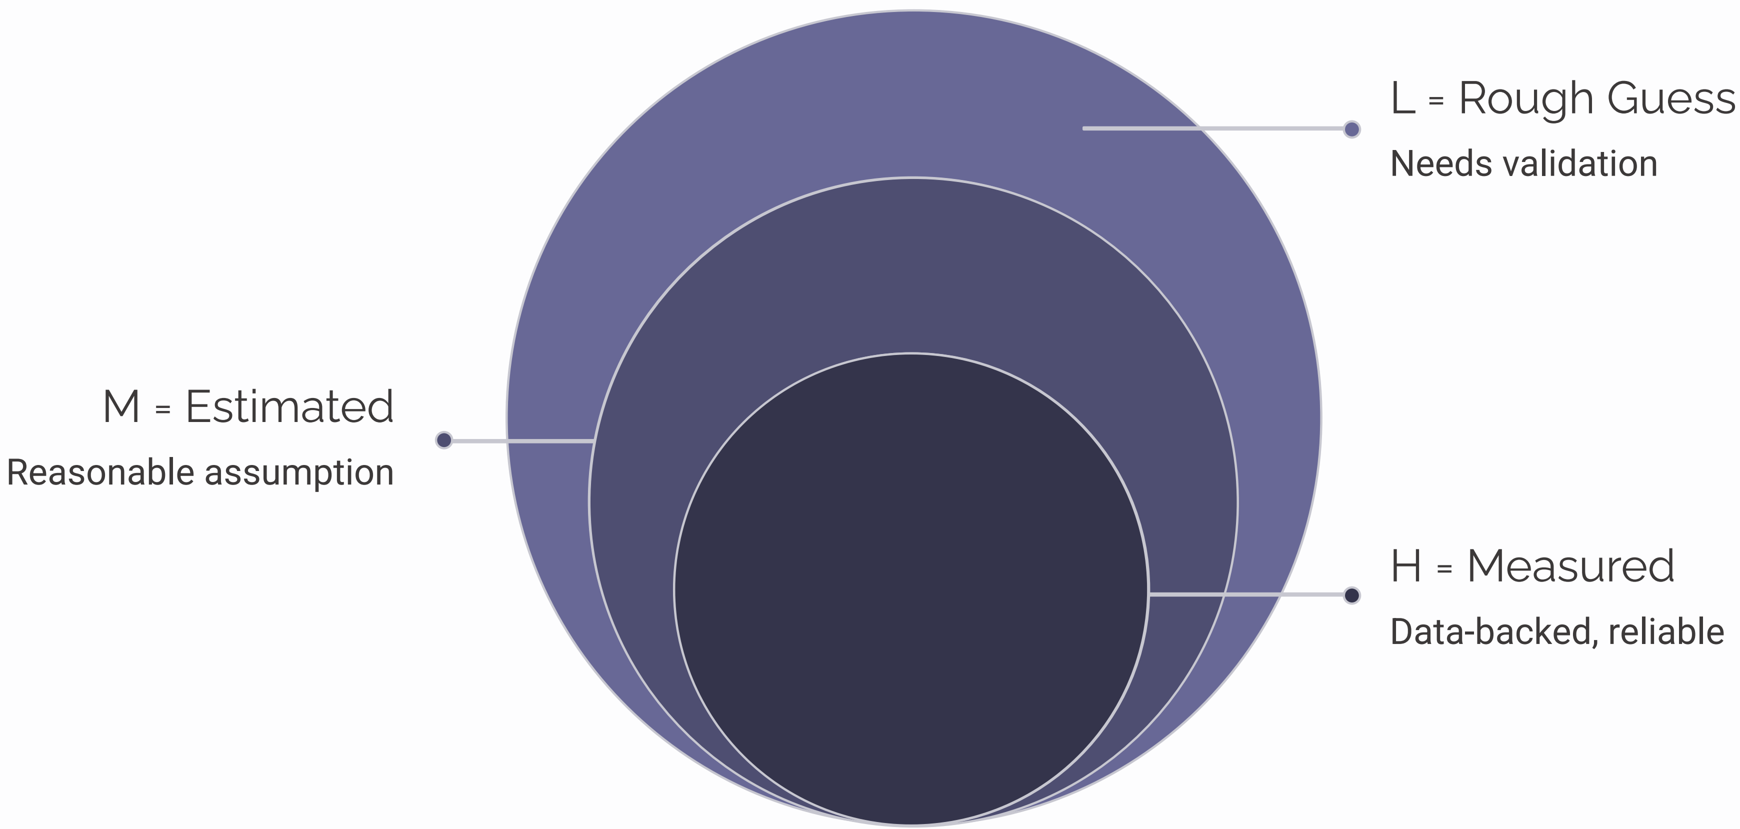

Confidence ratings guide how much weight leadership should place on each metric. Measured values (H) can be acted on immediately; rough guesses (L) should be flagged as requiring validation before decisions are made.

Section 4 — Showback View

First-pass cost mapped to teams, products, or namespaces. This is a **showback** — costs are made visible without billing teams directly.

Team / Namespace	Compute	Tooling	Ops Effort	Total	% of Total
Shared / Unallocated					
Total					100%

Key findings from the showback view — capture patterns, surprises, and early signals about where the cost levers are:

Section 5 — Key Findings & Questions

Use this space to capture findings that emerged during the session — numbers that surprised the group, patterns worth investigating, or early signals about where the levers are.

Key Findings

Numbers that surprised the group:

Patterns worth investigating:

Early signals about cost levers:

Questions and Notes

Open questions raised during the session:

Items to follow up on:

After This Session

Within **24 hours**, you will receive a recap email covering everything decided and produced in today's session.

Decisions Made

A full record of all allocation rule decisions agreed in the session

Unit Economics Model

Saved, versioned, and shared — the baseline metrics from today's work

Showback View

Saved, versioned, and shared — the first-pass cost-to-team mapping

Workshop 4 Confirmed

Date, time, and what to expect — the value and risk scorecard session

Workshop 4 preview: We look at the other side of the equation — not just what the platform costs, but what it *enables*. We will build a value and risk scorecard so cost is not the only lens leadership is looking through.

Glossary

Key terms used throughout this workshop and the broader cost model programme.

Term	Plain English
Unit economics	Cost per useful outcome — e.g. cost per app per month
Allocation rules	The agreed method for distributing shared costs
Showback	Making costs visible to teams without billing them
Chargeback	Billing teams directly for their platform usage
Numerator	The cost included in a unit metric calculation
Denominator	What you divide by — app count, tenant count, vCPU hours
Utilisation gap	The difference between capacity purchased and capacity used
Namespace	A Kubernetes construct for grouping and isolating workloads
Tenant	A team or product that uses a shared platform
Confidence rating	H = measured, M = estimated, L = rough guess
Assumptions Register	A versioned log of every scope and cost treatment decision
Tagging / labelling	Metadata applied to cloud resources to enable cost allocation
vCPU-hour	One virtual CPU running for one hour — the unit of compute measurement