


A stylized illustration of a workshop session. A group of five people are gathered around a table. One man stands in the background, while four others are seated. They are looking at laptops and documents on the table. The room has large windows, a potted plant, and framed artwork on the wall.

Workshop 4 — Attendee Materials

VALUE AND RISK SCORECARD

VERSION 1.0

This is your reference and working sheet for Workshop 4. Bring it to the session — either printed or open on your device. Some sections have space to capture measures, baselines, and target directions as we agree them in the room. Nothing needs to be completed in advance.

Session Date

[DATE]

Facilitator

[FACILITATOR]

What We Are Doing Today

Workshop 4 looks at the other side of the equation. We have built a cost model and a unit economics picture. Today we build a **value and risk scorecard** — what the platform enables, what it protects, and what would be at risk if cost decisions were made without considering it.

Value and Risk Scorecard

Eight to twelve agreed measures capturing platform value

Baseline Values

Confidence ratings attached to every baseline we agree

Target Directions

A clear direction for each measure — up, down, or stable

Value Narrative

A story connecting platform cost to platform value

Why This Matters

| Cost without value is an incomplete picture.

When leadership only sees cost, they optimise in one dimension — and often shift cost somewhere less visible, into incidents, audit failures, or developer friction. The scorecard gives leadership a balanced view and gives the platform team something they often lack: a way to articulate value in terms leadership recognises.

The Risk of Cost-Only Thinking

Cutting platform investment without a value lens shifts cost into hidden places — incidents, compliance failures, and engineering toil that never appear on a budget line.

What the Scorecard Provides

A structured, agreed set of measures that makes platform value visible to leadership in language they already use to make decisions.

What the Platform Team Gains

A defensible, evidence-based way to participate in cost conversations — not just as a cost centre, but as a value driver with measurable outcomes.

The Four Scorecard Categories

The scorecard is organised into four categories, each capturing a different dimension of platform value. Together they form a complete picture of what the platform enables and protects.



Delivery Performance

How effectively the platform enables teams to ship changes safely and quickly. Covers deployment frequency, lead time, and change failure rate.



Reliability

How stable the platform is and how quickly it recovers when things go wrong. Covers incident volume, MTTR, and platform availability.



Security and Compliance Effort

How much effort the platform requires and reduces for security and compliance. Covers patch time, policy coverage, and audit effort.



Developer Experience

How much friction developers experience when building and deploying. Covers onboarding time, time to first deployment, and toil volume.

-  Not every organisation tracks these measures formally. For each measure where direct data is not available, we agree a proxy – an imperfect but directionally useful substitute. A proxy with a Low confidence rating is more useful than no baseline at all.

Section 1 — Delivery Performance

AGREE MEASURES AND BASELINES IN THE SESSION

Delivery performance measures how effectively the platform enables teams to ship changes safely and quickly. Use the table below to capture agreed data sources, baselines, and confidence ratings during the session.

Measure	Data Source or Proxy	Baseline Value	Confidence	Target Direction
Deployment frequency			H / M / L	
Lead time for changes			H / M / L	
Change failure rate			H / M / L	

Proxy Options

Measure	Proxy
Deployment frequency	Release cadence per team per month
Lead time	Estimated time from code review to production
Change failure rate	Proportion of deployments followed by a hotfix or rollback



Key findings from this category: Capture notes here during the session.

Section 2 — Reliability

AGREE MEASURES AND BASELINES IN THE SESSION

Reliability measures how stable the platform is and how quickly it recovers when things go wrong. Stable, fast-recovering platforms reduce the hidden cost of incidents and protect engineering capacity.

Measure	Data Source or Proxy	Baseline Value	Confidence	Target Direction
Incident volume (per quarter)			H / M / L	
Time to restore service (MTTR)			H / M / L	
Platform availability			H / M / L	

Proxy Options

Measure	Proxy
Incident volume	On-call escalations per quarter
MTTR	Average incident duration from on-call logs
Availability	Rough uptime estimate from the platform team



Key findings from this category: Capture notes here during the session.

Section 3 — Security and Compliance Effort

AGREE MEASURES AND BASELINES IN THE SESSION

Security and compliance effort measures how much effort the platform requires and reduces for security and compliance activities. A well-built platform automates controls and reduces the manual burden of audit cycles.

Measure	Data Source or Proxy	Baseline Value	Confidence	Target Direction
Time to patch critical vulnerabilities			H / M / L	
Policy coverage (% of workloads)			H / M / L	
Audit evidence effort (days per cycle)			H / M / L	

Proxy Options

Measure	Proxy
Time to patch	Estimated average from recent CVE responses
Policy coverage	Rough percentage of workloads with enforced controls
Audit evidence effort	Estimated person-days per audit cycle



Key findings from this category: Capture notes here during the session.

Section 4 — Developer Experience

AGREE MEASURES AND BASELINES IN THE SESSION

Developer experience measures how much friction developers encounter when building and deploying. Toil, slow onboarding, and long paths to first deployment are invisible costs that compound over time across every team using the platform.

Measure	Data Source or Proxy	Baseline Value	Confidence	Target Direction
Onboarding time (new service or team)			H / M / L	
Time to first deployment			H / M / L	
Ticket and toil volume (per month)			H / M / L	

Proxy Options


Measure	Proxy
Onboarding time	Estimated from recent team onboarding experiences
Time to first deployment	Estimated from platform team experience
Ticket and toil volume	Support ticket count per month or per deployment

 **Key findings from this category:** Capture notes here during the session.

Section 5 — Scorecard Summary

Complete this together once all four categories are built. This is the consolidated view — the single artefact that will be saved, versioned, and shared after the session.

Category	Measure	Baseline	Confidence	Target Direction
Delivery Performance	Deployment frequency		H / M / L	
	Lead time for changes		H / M / L	
	Change failure rate		H / M / L	
Reliability	Incident volume		H / M / L	
	MTTR		H / M / L	
	Availability		H / M / L	
Security and Compliance	Time to patch		H / M / L	
	Policy coverage		H / M / L	
	Audit evidence effort		H / M / L	
Developer Experience	Onboarding time		H / M / L	
	Time to first deployment		H / M / L	
	Ticket and toil volume		H / M / L	



Section 6 — The Value Narrative

Capture the narrative that emerges from the scorecard. This is the story that connects platform cost to platform value — the text that will accompany the scorecard when it is shared with leadership.

1	2	3
What the Platform Is Enabling Well Record the areas where the scorecard shows strong performance or clear positive trends. These are the platform's demonstrable contributions to delivery, reliability, security, and developer productivity. <i>Notes:</i>	Where There Is Room to Improve Record the measures where baselines are weak, confidence is low, or target directions suggest significant gaps. These are honest areas for investment or attention — not weaknesses to hide. <i>Notes:</i>	What Would Be at Risk Record what would be at risk if cost were optimised without considering value. This is the most important part of the narrative — the consequences of cost decisions made without a value lens. <i>Notes:</i>

 This narrative section is the output that matters most to leadership. Spend time on it. Vague notes here will produce a weak case. Specific, evidence-linked statements will produce a strong one.

Questions and Notes

Use this space for any questions, observations, or parking-lot items that arise during the session and need follow-up.

After This Session

Within 24 hours you will receive a recap email covering everything agreed in the room. Here is what to expect next.



- ✓ Workshop 5 is the final session. Everything built across Workshops 1–4 comes together into a single, coherent platform investment case.

Glossary

Plain-English definitions for the key terms used throughout this workshop and the scorecard.

Term	Plain English
DORA	DevOps Research and Assessment – four metrics for software delivery performance
SPACE	A framework for developer productivity across five dimensions
Deployment frequency	How often code is successfully deployed to production
Lead time for changes	Time from code commit to running in production
Change failure rate	Percentage of deployments that cause a production failure
MTTR	Mean Time to Restore – average recovery time from incidents
Proxy	An imperfect but directionally useful substitute for a direct measurement
Policy coverage	The percentage of workloads governed by automated policy controls
Toil	Repetitive, manual, automatable work that does not add long-term value
Onboarding time	Time for a new service or team to reach first deployment
Confidence rating	H = measured, M = estimated, L = rough guess
Scorecard	A structured set of agreed measures with baseline and target values
Target direction	Whether a measure should increase, decrease, or stay the same