



Kubernetes Platform Economics


Workshop 2

Cost Taxonomy and Baseline Model

[DATE] | [FACILITATOR]

Where We Are

Track your progress across the full workshop series — you are here at Workshop 2.



What We Are Producing Today

By the end of this session we will have five concrete outputs ready for the team.

01

A Cost Taxonomy

Agreed and understood by everyone in the room.

02

A Baseline Model Skeleton

Known spend mapped in and accounted for.

03

Fixed vs. Variable Cost Split

A first pass at separating the two cost types.

04

A Clear View of Gaps

Each gap assigned an owner and a target date.

05

Spend Concentration Picture

An initial view of where the money is actually going.

Why Structure Matters

The difference between a taxonomy and no taxonomy is the difference between a trusted output and a contested one.

Without a Taxonomy

- Costs scattered across invoices, cloud bills, and people's heads
- IT and Finance disagree on what the numbers mean
- Output is contested and argued over
- No single source of truth to anchor decisions

With a Taxonomy

- Every cost has a home, a source, and an owner
- IT and Finance work from the same structure
- Output is auditable and trusted
- Decisions can be made with confidence

The Cost Taxonomy — Six Buckets

Every platform cost belongs in one of these six buckets. The structure is designed to be exhaustive, non-overlapping, and actionable.

1

Cloud Infrastructure

Compute, storage, and network costs

2

Managed Service Fees

Cloud provider charges for managed Kubernetes services

3

Platform Tooling

Observability, security, CI/CD, backup, and DR

4

Subscriptions

Platform licences and support agreements

5


Operating Effort

Platform team time — the largest hidden cost

6

Shared Services

IAM, shared network, shared logging



Bucket 1 — Cloud Infrastructure

The most visible layer of platform spend. Covers everything the cloud provider charges for raw compute, storage, and network resources.

Compute

- Virtual machines and node pools
- Reserved instances
- Spot usage

Storage

- Block storage
- Object storage
- Persistent volumes

Network

- Load balancers
- NAT gateways
- Data egress and peering

⚠ What to look for: Idle or oversized node pools · Egress costs suggesting inefficient data movement · Load balancer proliferation

Bucket 2 — Managed Service Fees

Fees charged by cloud providers for managed Kubernetes services. These are **separate from the compute cost** of running the nodes — a distinction that is frequently missed in cost reviews.

Examples

EKS

AWS cluster management fees per cluster per hour

AKS

Azure management fees for the managed control plane

Other Providers

Equivalent managed control plane charges from GKE and others

What to Look For

- How many clusters are currently incurring management fees?
- Are any clusters low-utilisation or candidates for retirement?
- Is the number of clusters growing without a corresponding growth in workloads?

Bucket 3 — Platform Tooling

Everything that makes the platforms operable, above and beyond raw infrastructure. This bucket is often fragmented across teams and billing accounts.



Observability

Monitoring, logging, alerting, and dashboards



Security

Scanning, policy enforcement, secrets management



CI/CD

Pipeline tooling and artefact management



Backup and DR

Backup services and disaster recovery tooling

⚠️ What to look for: The same capability running separately across multiple platforms · SaaS tooling costs that do not appear on the cloud bill · Tooling that was set up and never reviewed for value

Bucket 4 — Subscriptions

Platform licences, support agreements, and recurring software fees. These are **typically fixed costs** — they do not change with usage in the short term, which makes them both predictable and easy to overlook.



Predate Current Usage

Subscriptions that were set up before current platform usage patterns were established



Support Tier Fit

Support tiers that may be higher than current operational needs require



Licence Overage

Licences covering more capacity than is actively being used



Renewal Dates

Renewals approaching that represent an opportunity to renegotiate or exit

Bucket 5 — Operating Effort

The time your platform team spends running, maintaining, and improving the platforms. **Often the largest hidden cost**, and the one most commonly left out of cost reviews entirely.

What It Includes

- Upgrade and patching cycles
- Incident response and on-call
- Onboarding new teams and workloads
- Building and maintaining golden paths and standards

How We Capture It

A rough time split across activities is enough for this session.

We do not need timesheets. We need a reasonable estimate that reflects where the team's time actually goes.

 Precision is the enemy of progress here. A good estimate today is more valuable than a perfect number next month.



Bucket 6 — Shared Services

Costs that sit beneath the platforms and serve them, without being directly attributable to any single platform. These are real costs — they just require a different allocation approach.

Identity and Access Management

IAM infrastructure shared across all platforms and workloads

Shared Network Infrastructure

Core networking components that underpin multiple platforms

Shared Logging and Audit

Centralised logging or audit platforms serving the whole estate


Certificate Management

Centralised certificate issuance and renewal infrastructure

- ❏ Often the least complete bucket in the first pass. Capture what you can, note the gaps clearly, and move on. Completeness comes in later workshops.

Mapping Spend — How We Work Through It

We follow the same four-step approach for every cost bucket. The goal is a populated skeleton — not a finished model.



Speed matters here. A populated skeleton with honest confidence ratings is far more useful than a half-finished model where the team has spent an hour debating a single line item.

Fixed Versus Variable Costs

Understanding which costs are fixed and which are variable determines which levers are available — and how quickly you can pull them.

Fixed Costs

Do not change with usage in the short term. Require a **structural decision** to change.

- Subscriptions
- Support agreements
- Licence fees

Variable Costs

Scale with usage as demand changes. Require an **operational decision** to change.

- Compute consumption
- Storage usage
- Data egress

❓ Which lever is available to you right now — and how quickly can you pull it?

Data Gaps Review

Not every number will be available today. For each gap, the group agrees on one of three approaches before moving on.

Estimate	Wait	Exclude
Use a reasonable estimate. Assign a Low confidence rating and document the assumption.	Data is coming. Confirm a firm date for when it will be available and assign an owner.	Genuinely not available. Note the impact on overall confidence and move on.

Gap	Approach	Owner	Date
[ITEM]			
[ITEM]			
[ITEM]			



What the Baseline Is Telling Us

Populated live from session findings. These three summary points capture the most important signals from the baseline model as it stands today.

Where Spend Concentrates

[POPULATE IN SESSION]

Fixed Versus Variable Split

[POPULATE IN SESSION]

Key Gaps Still to Close

[POPULATE IN SESSION]

What We Have Agreed Today

A summary of everything the group has confirmed and committed to in this session.

- ☒ ~~Cost taxonomy agreed — six buckets, understood by the room~~
- ☒ ~~Baseline skeleton populated — known spend mapped in~~
- ☒ ~~Fixed versus variable split — first pass complete~~
- ☒ ~~Data gaps list updated — approach, owner, and date confirmed for each item~~
- ☒ ~~Assumptions Register updated with session findings~~
- ☐ Workshop 3 confirmed — [DATE] | [TIME]

☒ Strong session. The model is live and the gaps are visible. That is exactly where we need to be.



See You in Workshop 3

Allocation Rules and Unit Economics

[DATE] | [TIME] | [LOCATION / LINK]

Pre-Read

Coming your way before the session. Watch for it in your inbox.

Outstanding Data Items

Please action any open items from today's gaps list. The model updates as data comes in.