



Kubernetes Platform Economics

Workshop 3

Allocation Rules and Unit Economics

[DATE] | [FACILITATOR]

Where We Are

Track your progress across the full workshop series – you are here at Workshop 3.



What We Are Producing Today

By the end of this session, we will have produced five concrete outputs that form the foundation of your platform economics model.

1

Agreed Allocation Rules

Documented in the Assumptions Register

2

Unit Economics Model

Baseline values for three to five metrics

3

Showback View

Cost mapped to teams, products, or namespaces

4

Clear Narrative

What does the platform cost per outcome?

5

Foundation for Scenario Modelling

Ready for Workshop 5



Why Unit Economics — Not Just Total Cost

Total Cost

"We spend X on Kubernetes"

- Hard to act on
- Hard to compare
- Hard to defend

Unit Economics

"A production application costs Y per month"

- Leadership can act on this
- Finance can benchmark this
- Procurement can evaluate this

i The same spend. A completely different conversation. Total cost is a number. Unit economics is what that number **means**.

What Allocation Rules Are

Not all costs are created equal. Some are easy to attribute directly – others are genuinely shared across the estate.

Direct Costs

Straightforward to attribute

- A node pool running a specific workload
- A subscription that serves one platform

Shared Costs

Genuinely distributed across teams

- An observability stack serving three platforms
- A network platform serving the whole estate

- ❏ Allocation rules are the agreed method for distributing shared costs **fairly and consistently**. They do not need to be perfect. They need to be **agreed upon, documented, and applied consistently**.

Allocation — Shared Platform Costs

How do we distribute costs that serve multiple platforms? Three options are available — review the rule agreed in Workshop 1 against the actual numbers now visible.

1

Option A — Proportional by Usage

Most accurate — requires utilisation or tag data

2

Option B — Equal Split Across Platforms

Simpler — less precise

3

Option C — Shared Unallocated Bucket

Honest about uncertainty — refine later

⚠ We agreed in Workshop 1 to use: **[RULE FROM WORKSHOP 1]** — Does that still hold now that we can see the actual numbers?

Decision: [POPULATE IN SESSION]

Allocation — Platform Tooling

How do we allocate tooling that serves multiple platforms? This covers observability, security, and CI/CD infrastructure.

Option A

Allocate by cluster or node count per platform

Option B

Allocate equally across platforms

Option C

Hold as shared — do not allocate to platforms

 Where tooling duplication exists, model it as a **separate line** rather than allocating it. This keeps the model honest and surfaces duplication as a visible cost.

Decision: [POPULATE IN SESSION]

Allocation — Operating Effort

How do we allocate platform team time across platforms and workloads? People cost is often the largest single line item — the allocation method matters.

Option A — Allocate by Estimated Time Spent

Allocate by estimated time spent per platform. Most reflective of actual effort — requires time tracking or estimation.

Option B — Equal Split Across Platforms

Allocate equally across platforms. Simple and defensible — may not reflect reality.

Option C — Platform-Level Cost

Hold as a platform-level cost — do not allocate to individual workloads. Keeps complexity low — limits showback granularity.

Decision: [POPULATE IN SESSION]

Minimum Tagging Standard

To allocate costs by namespace or team, we need a minimum tagging standard in place. The table below should be populated live during the session to capture current state and identify gaps.

 Without consistent tags or labels, showback is an estimate at best. Closing tagging gaps is a prerequisite for reliable unit economics.

Tag / Label	What It Identifies	Coverage
[TAG 1]	[POPULATE IN SESSION]	[POPULATE IN SESSION]
[TAG 2]	[POPULATE IN SESSION]	[POPULATE IN SESSION]
[TAG 3]	[POPULATE IN SESSION]	[POPULATE IN SESSION]

What is the minimum we need for showback to work?

Gaps to close: [POPULATE IN SESSION]

The Unit Metrics We Are Building

In Workshop 1, we shortlisted the metrics below. For each one, we will define the numerator, denominator, data sources, and confidence level during this session.

01 [Metric 1] Numerator: Which costs are included Denominator: What we divide by Data sources: Where the numbers come from Confidence: H / M / L	02 [Metric 2] Numerator: Which costs are included Denominator: What we divide by Data sources: Where the numbers come from Confidence: H / M / L	03 [Metric 3] Numerator: Which costs are included Denominator: What we divide by Data sources: Where the numbers come from Confidence: H / M / L
04 [Metric 4] Numerator: Which costs are included Denominator: What we divide by Data sources: Where the numbers come from Confidence: H / M / L	05 [Metric 5] Numerator: Which costs are included Denominator: What we divide by Data sources: Where the numbers come from Confidence: H / M / L	

Metric — Cost Per Production App Per Month

What it tells you: What it costs to run one production application on this platform. This is the headline metric for platform value conversations with leadership.

Definition

Numerator: Total platform cost allocated to production environments

Denominator: Number of production applications

Populate in Session

Data sources: [POPULATE IN SESSION]

Baseline value: [POPULATE IN SESSION]

Confidence: H / M / L

What does this number enable? [POPULATE IN SESSION]

Metric — Cost Per Tenant or Namespace Per Month

What it tells you: What it costs to serve one team or product on the platform. This metric drives showback conversations and surfaces disproportionate consumers.

Definition

Numerator: Total platform cost allocated to a team or namespace

Denominator: Number of active tenants or namespaces

Populate in Session

Data sources: [POPULATE IN SESSION]

Baseline value: [POPULATE IN SESSION]

Confidence: H / M / L

What does this number enable? [POPULATE IN SESSION]

Metric — Cost Per Environment

What it tells you: How much of the total platform cost is dedicated to production versus non-production environments.

Production Cost [POPULATE] — [%] of total	Non-Production Cost [POPULATE] — [%] of total
--	--

Data sources: [POPULATE IN SESSION] | **Confidence:** H / M / L

 **What does a high non-production share signal?** Cluster sprawl. Over-provisioning in dev and test. Environments that run continuously when they could be ephemeral.

Metric — Effective Cost Per vCPU-Hour Used

What it tells you: How efficiently you are using the capacity you pay for. The gap between purchased and used is often where the most straightforward cost improvements sit.

[?]

vCPU Hours Purchased
Allocated capacity

[?]

vCPU Hours Requested
Scheduled by workloads

[?]

vCPU Hours Actually Used
Real consumption

[?]%

Utilisation Gap
Purchased vs. used

Effective cost per vCPU-hour used: [POPULATE] | **Confidence:** H / M / L

Metric — Ops Cost Per Cluster Per Month

What it tells you: The operational overhead of each cluster in people cost. A high ops cost per cluster signals cluster sprawl — clusters that are disproportionately expensive to operate relative to the workloads they run.

Definition

Numerator: Operating effort cost allocated to clusters

Denominator: Number of clusters

Populate in Session

Data sources: [POPULATE IN SESSION]

Baseline value: [POPULATE IN SESSION]

Confidence: H / M / L

What does this number enable? [POPULATE IN SESSION]

 A high ops cost per cluster often signals **cluster sprawl** — clusters that are disproportionately expensive to operate relative to the workloads they run.

Unit Economics Summary

Summary of all metrics populated during this session. Baseline values and confidence levels to be confirmed live.

Metric	Baseline Value	Confidence
[Metric 1]	[VALUE]	H / M / L
[Metric 2]	[VALUE]	H / M / L
[Metric 3]	[VALUE]	H / M / L
[Metric 4]	[VALUE]	H / M / L
[Metric 5]	[VALUE]	H / M / L

 **Key findings:** [POPULATE IN SESSION]

Showback View

Making costs visible — not billing anyone. This first-pass showback maps platform spend to teams and namespaces using the allocation rules agreed today.

Team / Namespace	Compute	Tooling	Ops Effort	Total	% of Total
[TEAM / NS 1]	[POPULATE]	[POPULATE]	[POPULATE]	[POPULATE]	[POPULATE]
[TEAM / NS 2]	[POPULATE]	[POPULATE]	[POPULATE]	[POPULATE]	[POPULATE]
[TEAM / NS 3]	[POPULATE]	[POPULATE]	[POPULATE]	[POPULATE]	[POPULATE]
[TEAM / NS 4]	[POPULATE]	[POPULATE]	[POPULATE]	[POPULATE]	[POPULATE]
Shared / Unallocated	[POPULATE]	[POPULATE]	[POPULATE]	[POPULATE]	[POPULATE]
Total	[POPULATE]	[POPULATE]	[POPULATE]	[POPULATE]	100%

What We Have Agreed Today

A summary of everything confirmed and completed in this session.

- ✓ Allocation rules agreed — documented in the Assumptions Register
- ✓ Unit economics model populated — baseline values confirmed
- ✓ Showback view built — first pass complete
- ✓ Assumptions Register updated
- ☐ Workshop 4 confirmed — **[DATE]** | **[TIME]**

See You in Workshop 4

Value and Risk Scorecard

[DATE] | [TIME] | [LOCATION / LINK]

Pre-read coming your way before the session.

 Workshop 4 looks at the other side of the equation — **what value the platform is actually delivering**. We move from cost to outcomes.