

Kubernetes Platform Economics

Workshop 4 — Value and Risk Scorecard

[DATE] | [FACILITATOR]



Where We Are

Track your progress through the full workshop series — you are here at Workshop 4.



What We Are Producing Today

By the end of this session we will have five concrete outputs ready to carry forward into Workshop 5.

1

Value and Risk Scorecard

Eight to twelve agreed measures that define platform performance.

2

Baseline Values

Real data where available, proxies where not – every measure grounded.

3

Target Directions

Where each measure should move and why – agreed as a group.

4

A Narrative

What the platform enables and what it protects – in plain language.

5

Foundation for Workshop 5

A balanced basis for scenario evaluation in the next session.

Cost Without Value Is an Incomplete Picture

Optimising cost without understanding value shifts spend somewhere less visible — it does not eliminate it.

Platform That Looks Expensive

May be delivering significant returns in:

- Delivery speed
- Reliability
- Security posture
- Developer productivity

Platform That Looks Cheap

May be hiding cost in:

- Incidents and outages
- Audit failures
- Developer friction
- Manual compliance work

⚠ Optimising cost without value reduces spend in one place and shifts it somewhere less visible.



The Frameworks We Use

Both frameworks are industry standards — not vendor-specific — and together they give us a complete picture of platform performance and developer productivity.

DORA

DevOps Research and Assessment

Four metrics for software delivery performance:

- Deployment frequency
- Lead time for changes
- Change failure rate
- Time to restore service

SPACE

Developer Productivity

Five dimensions for thinking about developer productivity:

- Satisfaction and wellbeing
- Performance
- Activity
- Communication and collaboration
- Efficiency and flow

The Four Scorecard Categories



Delivery Performance

How effectively the platform enables teams to ship changes safely and quickly.



Reliability

How stable the platform is and how quickly it recovers when things go wrong.



Security and Compliance Effort

How much effort the platform requires – and reduces – for security and compliance obligations.



Developer Experience

How much friction developers experience when building and deploying on the platform.



A Note on Proxies

Not every organisation tracks DORA metrics formally. That is fine. For each measure we will agree a proxy if direct measurement is not available.

Measure	Direct	Proxy if not tracked	Confidence
Deployment frequency	Tracked in CI/CD	Release cadence per month	High
Lead time	Tracked in Git	Estimated by platform team	Medium
MTTR	Tracked in ops	Average duration from on-call logs	Medium
Onboarding time	Not tracked	Estimated from recent examples	Low

 A proxy with **Low** confidence is better than no baseline at all.

Category 1 — Delivery Performance

LIVE POPULATION

How effectively does the platform enable teams to ship changes safely and quickly? Populate the table below during the session.

Measure	Data Source / Proxy	Baseline	Confidence
Deployment frequency			
Lead time for changes			
Change failure rate			

 Use CI/CD pipeline data, Git history, or platform team estimates as your primary sources.

Category 2 — Reliability

LIVE POPULATION

How stable is the platform — and how quickly does it recover when things go wrong? Populate the table below during the session.

Measure	Data Source / Proxy	Baseline	Confidence
Incident volume (per quarter)			
Time to restore service (MTTR)			
Platform availability			

 Draw from on-call logs, incident management tools, and uptime monitoring dashboards.

Category 3 — Security and Compliance Effort

LIVE POPULATION

How much effort does the platform require to meet security and compliance obligations — and how much does it reduce that effort? Populate the table below during the session.

Measure	Data Source / Proxy	Baseline	Confidence
Time to patch critical vulnerabilities			
Policy coverage (% of workloads)			
Audit evidence effort (days per cycle)			

 Use vulnerability scanner reports, policy engine dashboards, and audit team estimates.

Category 4 — Developer Experience

LIVE POPULATION

How much friction do developers experience when building and deploying on the platform? Populate the table below during the session.

Measure	Data Source / Proxy	Baseline	Confidence
Onboarding time (new service or team)			
Time to first deployment			
Ticket and toil volume per month			

 Use developer surveys, ticket system data, and estimates from recent onboarding examples.

Scorecard Summary

LIVE POPULATION

All twelve measures in one view — populate baselines, confidence ratings, and target directions together as a group.

Category	Measure	Baseline	Confidence	Direction
Delivery Performance	Deployment frequency			
Delivery Performance	Lead time			
Delivery Performance	Change failure rate			
Reliability	Incident volume			
Reliability	MTTR			
Reliability	Availability			
Security and Compliance	Time to patch			
Security and Compliance	Policy coverage			
Security and Compliance	Audit effort			
Developer Experience	Onboarding time			
Developer Experience	Time to first deploy			
Developer Experience	Ticket and toil volume			



Agreeing Target Directions

For each measure — where should it move, and over what timeframe? We are agreeing **directions**, not precise targets.

- ☐ Should this increase, decrease, or stay the same?
Agree the direction of travel for each measure as a group before moving to magnitude.
- ☐ What is a realistic improvement over six to twelve months?
Anchor expectations to what is achievable given current team capacity and platform maturity.
- ☐ What would achieving that improvement be worth?
Quantify the value — in cost reduction, risk reduction, or team capacity freed up.

What the Scorecard Tells Us

LIVE POPULATION

Populated live from session findings — capture the group's conclusions in each area below.

What the platform is enabling well

[POPULATE IN SESSION]

Where there is room to improve

[POPULATE IN SESSION]

What would be at risk if cost were optimised without considering value

[POPULATE IN SESSION]

What We Have Agreed Today

A summary of everything confirmed and completed in this session – ready to carry forward.

- ☒ ~~Value and risk scorecard built — [NUMBER] measures agreed~~
- ☒ ~~Baselines confirmed — with confidence ratings throughout~~
- ☒ ~~Target directions agreed for each measure~~
- ☒ ~~Scorecard saved and versioned~~
- ☒ ~~Assumptions Register updated~~
- ☐ Workshop 5 confirmed — [DATE] | [TIME]

☒ Great work. The scorecard is the foundation for everything in Workshop 5.

See You in Workshop 5

Scenarios, Recommendation, and 90-Day Plan

[DATE] | [TIME] | [LOCATION / LINK]

Pre-read coming your way before the session.

Workshop 5 brings everything together — cost model, unit economics, and value scorecard — and turns it into a decision and a plan.

Cost Model

From Workshops 1 & 2

Unit Economics

From Workshop 3

Value Scorecard

From Workshop 4 — today

Decision and Plan

Workshop 5 output